

# Chapter 7

## Entity Relationship Model

Slides by Silbershatz, Modifications by Rogers and Brown

# Entity-Relationship Model

- Design Process
- Modeling
- Constraints
- E-R Diagram
- Design Issues
- Weak Entity Sets
- Extended E-R Features
- Design of the Bank Database
- Reduction to Relation Schemas
- Database Design
- UML

# Modeling

- A database can be modeled as:
  - A collection of entities,
  - Relationship among entities.
- An **entity** is an object that exists and is distinguishable from other objects.
  - Example: specific person, company, event, plant
- Entities have **attributes**
  - Example: people have names and addresses
- An **entity set** is a set of entities of the same type that share the same properties.
  - Example: set of all persons, companies, trees, holidays

# Entity Sets instructor and student

Instructor_ID	Instructor_Name
76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

Student_ID	Student_Name
98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

# Relationship Sets

- A relationship is an association among several entities

Example:

44553 (Peltier)	Advisor	22222 (Einstein)
student entity	relationship set	instructor entity

- A relationship set is a mathematical relation among  $n \geq 2$  entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where  $(e_1, e_2, \dots, e_n)$  is a relationship

- Example :

$(44553, 22222) \in \text{advisor}$

# Relationship Set advisor

Instructor_ID	Instructor_Name
76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

Student_ID	Student_Name
98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier


student

# Relationship Sets (cont)

- At attribute can also be a property of a relationship set
- For instance, the advisor relationship set between entity sets instructor and student may have the attribute date which tracks when the student started being associated with the advisor

Instructor_ID	Instructor_Name
76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor



Student_ID	Student_Name
98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

# Degree of a Relationship Set

- **Binary relationship**
  - Involve two entity sets (or degree two).
  - Most relationship sets in a database system are binary.
- Relationships between more than two entity sets are rare. Most relationships are binary. (More on this later.)
  - Example: *students* work on research *projects* under the guidance of an *instructor*.
  - Relationship *proj\_guide* is a ternary relationship between *instructor*, *student*, and *project*




# Attributes

- An entity is represented by a set of attributes, that is descriptive properties possessed by all members of an entity set.
- Example:  
instructor = (ID, name, street, city, salary )  
course= (course\_id, title, credits)
- **Domain** – the set of permitted values for each attribute
- Attribute types:
  - **Simple** and **composite** attributes.
  - **Single-valued** and **multivalued** attributes
    - Example: multivalued attribute: phone\_numbers
  - **Derived** attributes
    - Can be computed from other attributes
    - Example: age, given date\_of\_birth

# Composite Attributes

- Composite Attributes




# Mapping Cardinality Constraints


- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
  - One to one
  - One to many
  - Many to one
  - Many to many

# Mapping Cardinalities

- **Note:** Some elements in A and B may not be mapped to any elements in the other set




(a)  
One to One




(b)  
One to Many

# Mapping Cardinalities

- **Note:** Some elements in A and B may not be mapped to any elements in the other set



(a)  
Many to One



(b)  
Many to Many

# Keys

- A **super key** of an entity set is a set of one or more attributes whose values uniquely determine each entity.
- A **candidate key** of an entity set is a minimal super key
  - *ID* is candidate key of *instructor*
  - *course\_id* is candidate key of *course*
- Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.

# Keys for Relationship Sets

- The combination of primary keys of the participating entity sets forms a super key of a relationship set.
  - (s\_id, i\_id) is the super key of advisor
  - NOTE: this means a pair of entity sets can have **at most** one relationship in a particular relationship set.
    - Example: if we wish to track multiple meeting dates between a student and her advisor, we cannot assume a relationship for each meeting. We can use a multivalued attribute though
- Must consider the mapping cardinality of the relationship set when deciding what are the candidate keys
- Need to consider semantics of relationship set in selecting the primary key in case of more than one candidate key

# Redundant Attributes

- Suppose we have entity sets
  - instructor, with attributes including dept\_name
  - departmentand a relationship
  - inst\_dept relating instructor and department
- Attribute dept\_name in entity instructor is redundant since there is an explicit relationship inst\_dept which relates instructors to departments
  - The attribute replicates information present in the relationship, and should be removed from instructor
  - BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see.



# E-R Diagrams




- Rectangles represent entity sets
- Diamonds represent relationship sets
- Attributes listed inside entity rectangle
- Underline indicates primary key attributes

# Entity with Composite, Multivalued and Derived Attributes

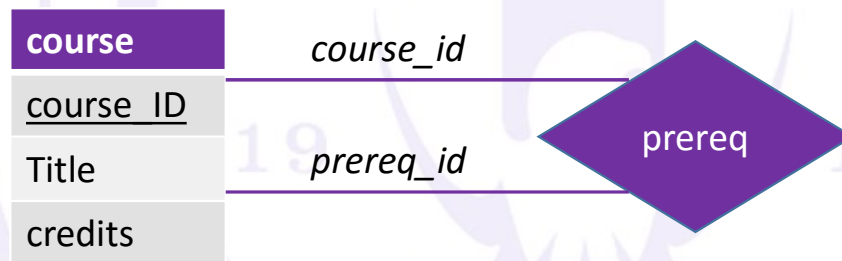
instructor	
ID	
name	
first_name	
middle_initial	
last_name	
address	
street	
street_number	
street_name	
apt_number	
city	
state	
zip	
{ phone_number }	
date_of_birth	
age ( )	

# Relationship Sets with Attributes



# Roles

- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a “role” in the relationship
- The labels “course\_id” and “prereq\_id” are called **roles**



# Cardinality Constraints

- We express cardinality constraints by drawing either a directed line ( $\rightarrow$ ), signifying “one,” or an undirected line ( $-$ ), signifying “many,” between the relationship set and the entity set.
- One-to-one relationship:
- A student is associated with at most one *instructor* via the relationship *advisor*
- A *student* is associated with at most one *department* via *stud\_dept*

# One-to-One Relationship

- One-to-one relationship between an instructor and a student
  - An instructor is associated with at most one student via advisor
  - A student is associated with at most one instructor via advisor



# One-to-Many Relationship

- One-to-many relationship between an instructor and a student
  - An instructor is associated with several (including 0) students via advisor
  - A student is associated with at most one instructor via advisor



# Many-to-One Relationship

- Many-to-one relationship between an instructor and a student
  - An instructor is associated with at most one students via advisor
  - A student is associated with several (including 0) instructor via advisor





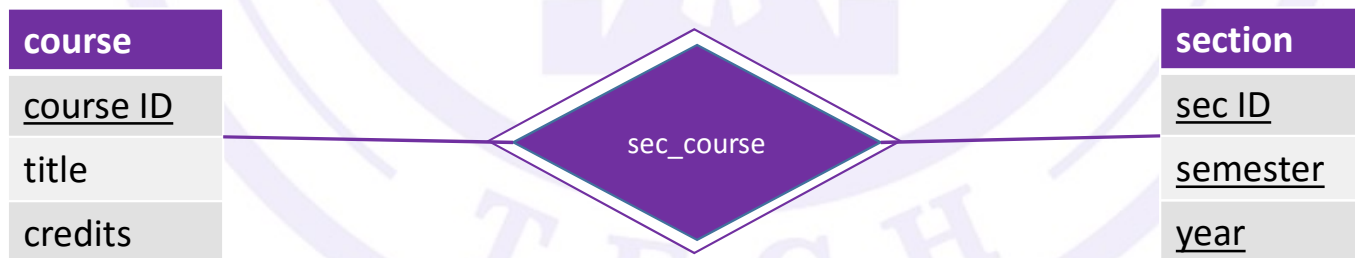
# Many-to-Many Relationship

- Many-to-many relationship between an instructor and a student
  - An instructor is associated with several (including 0) students via advisor
  - A student is associated with several (including 0) instructor via advisor



# Participation of an Entity Set in a Relationship Set

- Total participation (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
  - E.g., participation of section in sec\_course is total
    - every section must have an associated course
- Partial participation: some entities may not participate in any relationship in the relationship set
  - Example: participation of instructor in advisor is partial

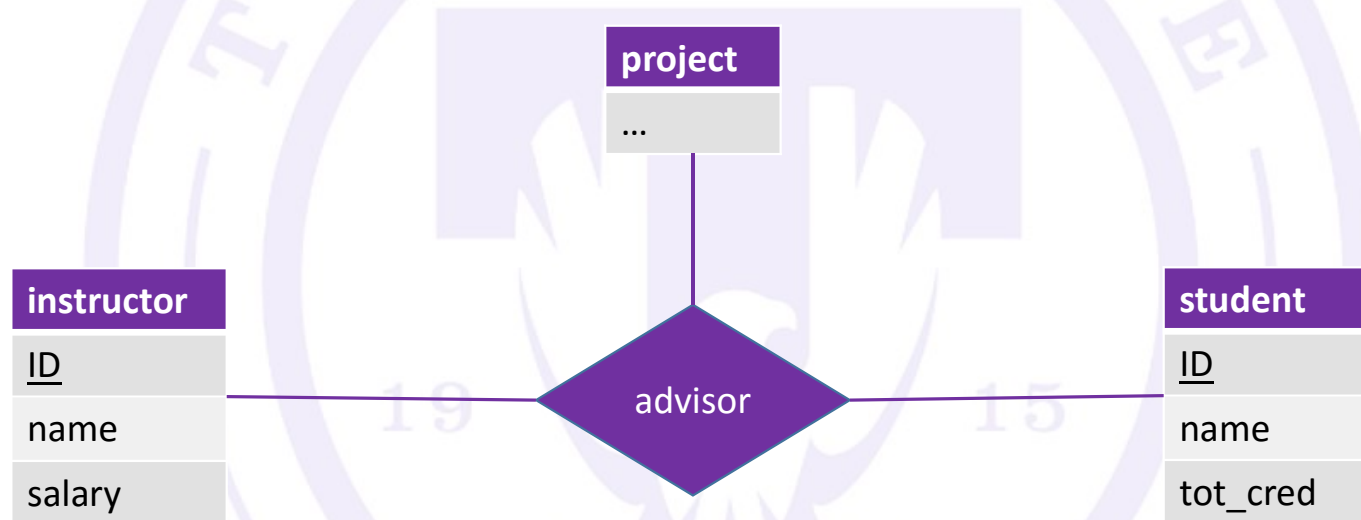


# Alternative Notation for Cardinality Limits

- Cardinality limits can also express participation constraints



# E-R Diagram with a Ternary Relationship



# Cardinality Constraints on Ternary Relationship

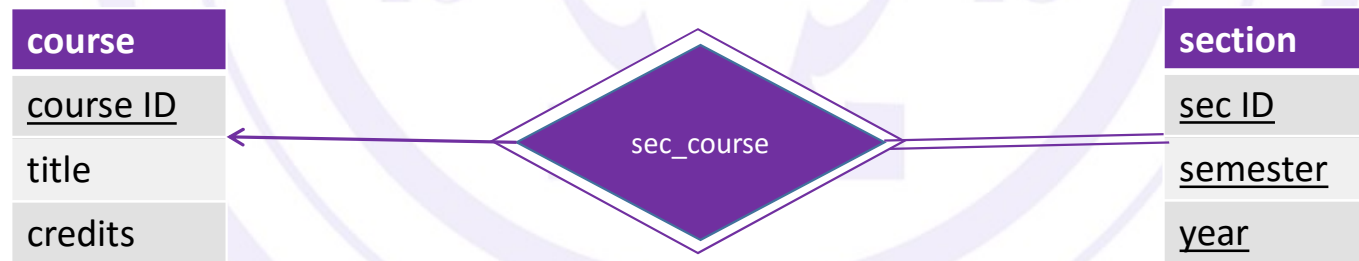
- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
- E.g., an arrow from proj\_guide to instructor indicates each student has at most one guide for a project
- If there is more than one arrow, there are two ways of defining the meaning.
  - E.g., a ternary relationship R between A, B and C with arrows to B and C could mean
    1. each A entity is associated with a unique entity from B and C or
    2. each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B
  - Each alternative has been used in different formalisms
  - To avoid confusion we outlaw more than one arrow

# Weak Entity Sets

- An entity set that does not have a primary key is referred to as a **weak entity set**.
- The existence of a weak entity set depends on the existence of a **identifying entity set**
  - It must relate to the identifying entity set via a total, one-to-many relationship set from the identifying to the weak entity set
  - **Identifying relationship** depicted using a double diamond
- The **discriminator** (or partial key) of a weak entity set is the set of attributes that distinguishes among all the entities of a weak entity set.
- The primary key of a weak entity set is formed by the primary key of the strong entity set on which the weak entity set is existence dependent, plus the weak entity set's discriminator.

# Weak Entity Sets

- We underline the discriminator of a weak entity set with a dashed line.
- We put the identifying relationship of a weak entity in a double diamond.
- Primary key for *section* – (*course\_id*, *sec\_id*, *semester*, *year*)

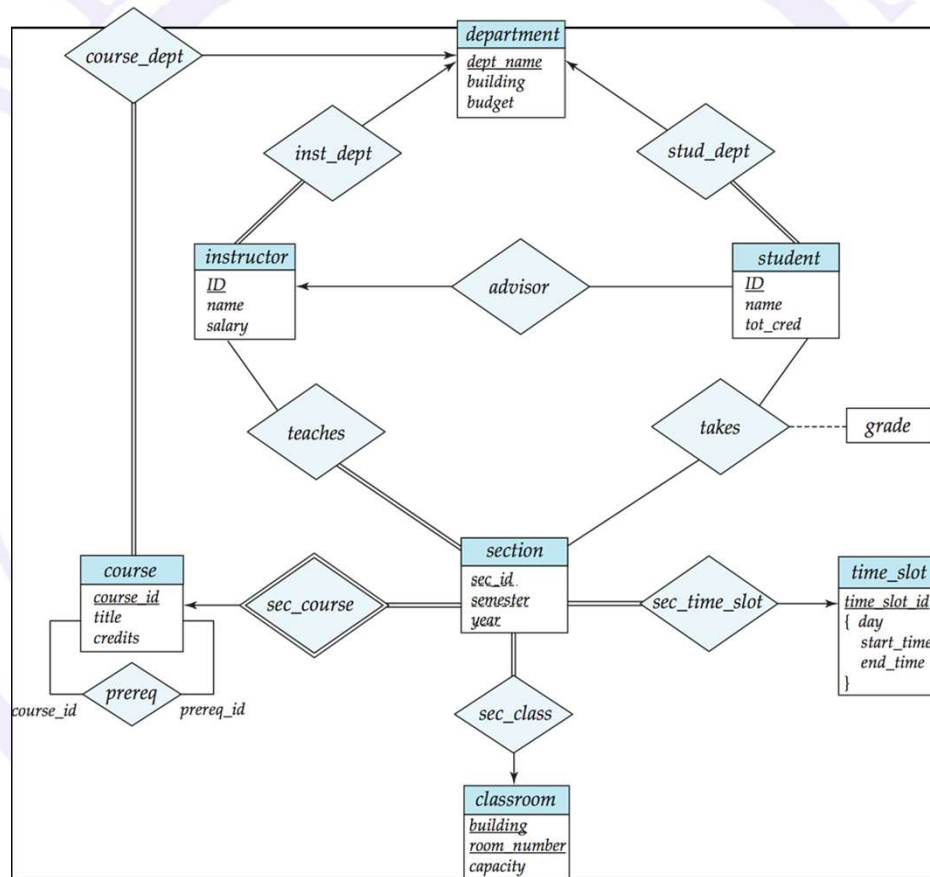


# Weak Entity Sets

- Note: the primary key of the strong entity set is not explicitly stored with the weak entity set, since it is implicit in the identifying relationship.
- If `course_id` were explicitly stored, `section` could be made a strong entity, but then the relationship between `section` and `course` would be duplicated by an implicit relationship defined by the attribute `course_id` common to `course` and `section`



# E-R Diagram for a University Enterprise

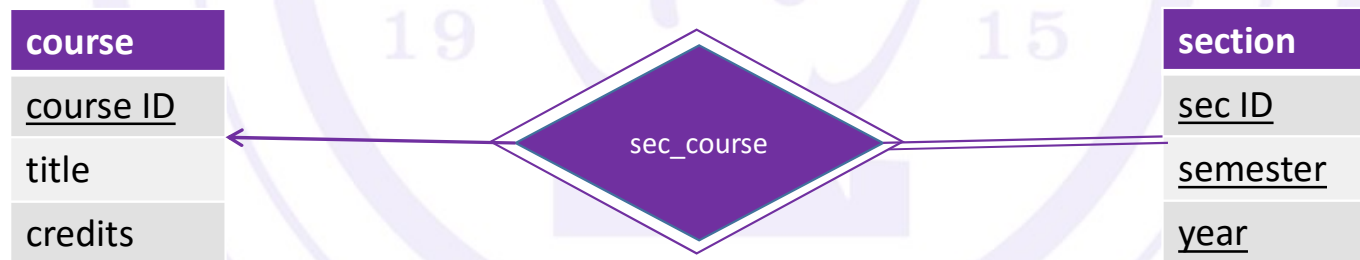


# Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as **relation schemas** that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.

# Representing Entity Sets with Simple Attributes

- A strong entity set reduces to a schema with the same attributes *student*(ID, name, tot\_cred)
- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set *section* ( course id, sec id, sem, year )



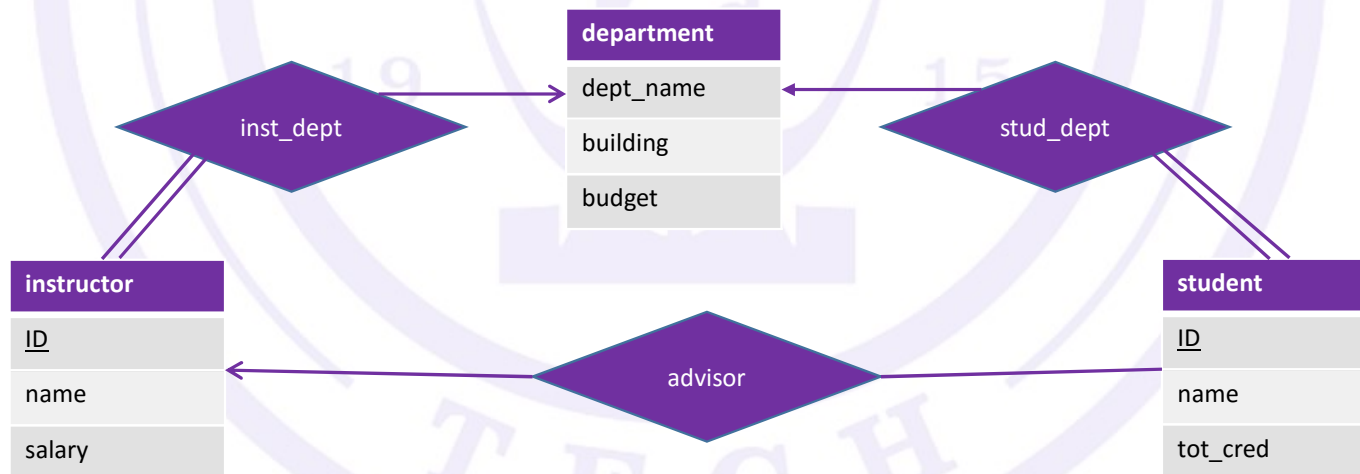
# Representing Relationship Sets

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*  
*advisor* = (*s\_id*, *i\_id*)



# Redundancy of Schemas

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side
- Example: Instead of creating a schema for relationship set inst\_dept, add an attribute dept\_name to the schema arising from entity set instructor



# Redundancy of Schemas

- For one-to-one relationship sets, either side can be chosen to act as the “many” side
  - That is, extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is partial on the “many” side, replacing a schema by an extra attribute in the schema corresponding to the “many” side could result in null values
- The schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.
  - Example: The section schema already contains the attributes that would appear in the sec\_course schema

# Composite and Multivalued Attributes

instructor	
ID	
name	
first_name	
middle_initial	
last_name	
address	
street	
street_number	
street_name	
apt_number	
city	
state	
zip	
{ phone_number }	
date_of_birth	
age ( )	

- Composite attributes are flattened out by creating a separate attribute for each component attribute
  - Example: given entity set instructor with composite attribute name with component attributes first\_name and last\_name the schema corresponding to the entity set has two attributes name\_first\_name and name\_last\_name
    - Prefix omitted if there is no ambiguity
- Ignoring multivalued attributes, extended instructor schema is
  - instructor(ID, first\_name, middle\_initial, last\_name, street\_number, street\_name, apt\_number, city, state, zip\_code, date\_of\_birth)



# Composite and Multivalued Attributes

- A multivalued attribute M of an entity E is represented by a separate schema EM
  - Schema EM has attributes corresponding to the primary key of E and an attribute corresponding to multivalued attribute M
  - Example: Multivalued attribute phone\_number of instructor is represented by a schema:  
inst\_phone= ( ID, phone\_number)
  - Each value of the multivalued attribute maps to a separate tuple of the relation on schema EM
    - For example, an instructor entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:  
(22222, 456-7890) and (22222, 123-4567)



# Multivalued Attributes

- Special case: entity `time_slot` has only one attribute other than the primary-key attribute, and that attribute is multivalued
  - Optimization: Don't create the relation corresponding to the entity, just create the one corresponding to the multivalued attribute
  - `time_slot(time_slot_id, day, start_time, end_time)`
  - Caveat: `time_slot` attribute of `section` (from `sec_time_slot`) cannot be a foreign key due to this optimization

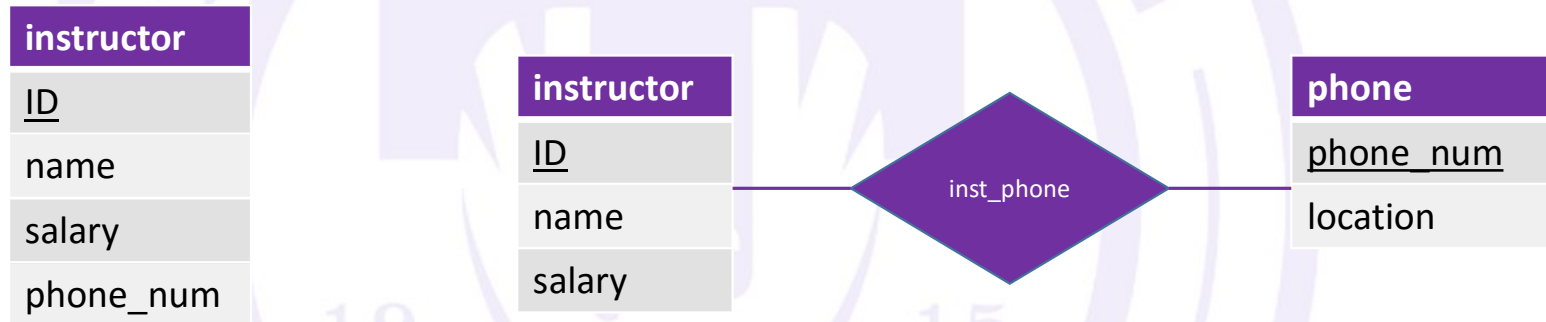
section
<u>sec ID</u>
<u>semester</u>
<u>year</u>

sec\_time\_slot

time_slot
time_slot_ID
{ day start_time end_time }

# Design Issues

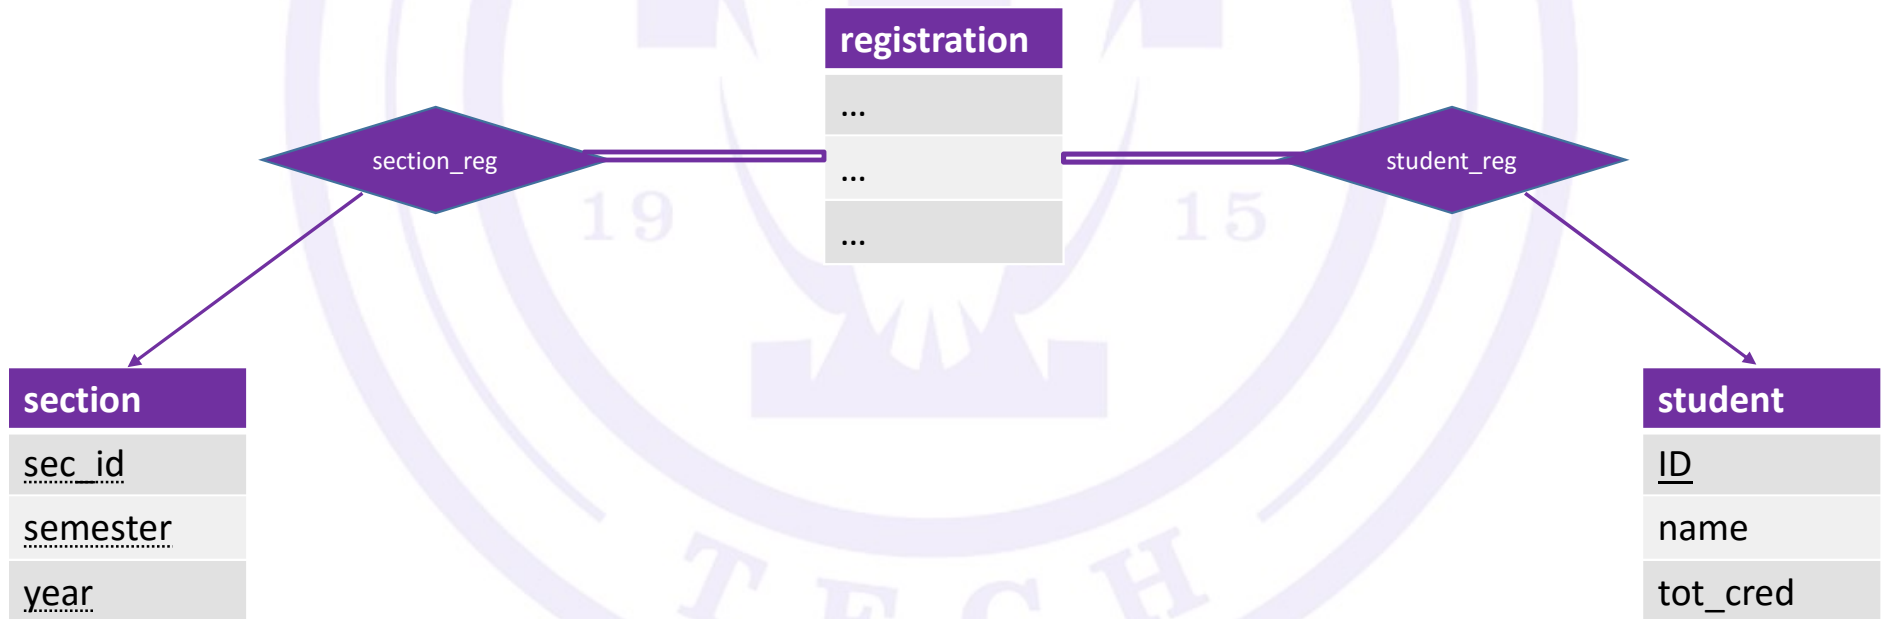
- Use of entity sets v attributes



- Use of phone as an entity allows extra information about phone numbers (plus multiple phone numbers)

# Design Issues

- Use of entity sets vs. relationship sets  
Possible guideline is to designate a relationship set to describe an action that occurs between entities



# Design Issues

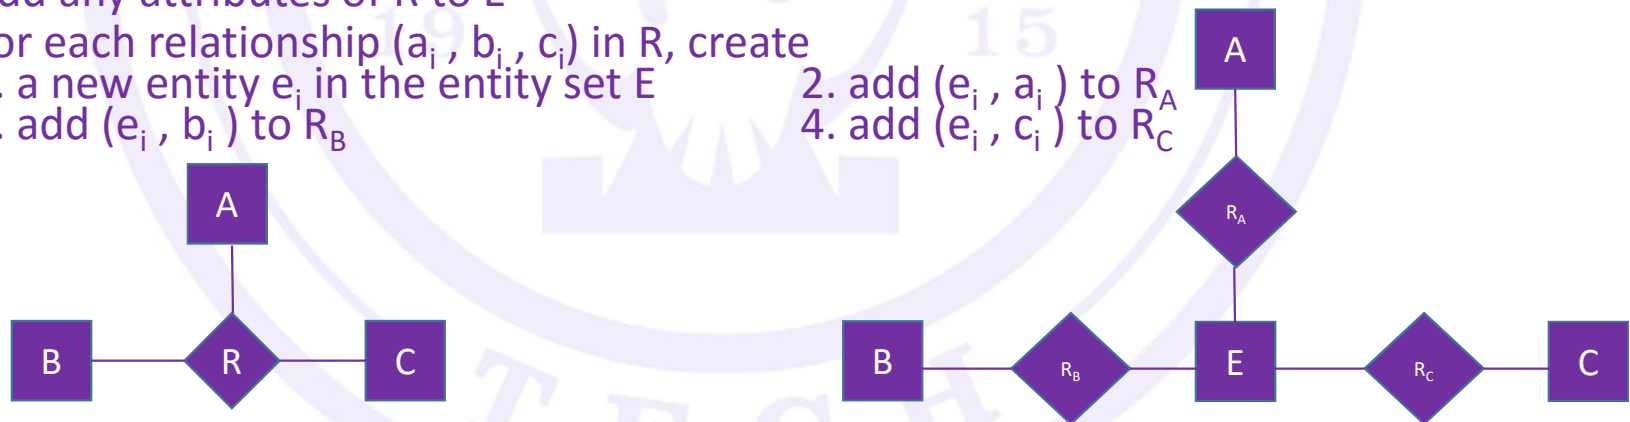
- Binary versus  $n$ -ary relationship sets
  - Although it is possible to replace any nonbinary ( $n$ -ary, for  $n > 2$ ) relationship set by a number of distinct binary relationship sets, a  $n$ -ary relationship set shows more clearly that several entities participate in a single relationship.
- Placement of relationship attributes  
e.g., attribute date as attribute of advisor or as attribute of student

# Binary v Non-Binary Relationships

- Some relationships that appear to be non-binary may be better represented using binary relationships
  - E.g., A ternary relationship parents, relating a child to his/her father and mother, is best replaced by two binary relationships, father and mother
    - Using two binary relationships allows partial information (e.g., only mother being know)
  - But there are some relationships that are naturally non-binary
    - Example: proj\_guide

# Converting Non-Binary Relationships to Binary Form

- In general, any non-binary relationship can be represented using binary relationships by creating an artificial entity set.
  - Replace R between entity sets A, B and C by an entity set E, and three relationship sets:
- 1.  $R_A$ , relating E and A
  2.  $R_B$ , relating E and B
  3.  $R_C$ , relating E and C
  - Create a special identifying attribute for E
  - Add any attributes of R to E
  - For each relationship  $(a_i, b_i, c_i)$  in R, create
    1. a new entity  $e_i$  in the entity set E
    2. add  $(e_i, a_i)$  to  $R_A$
    3. add  $(e_i, b_i)$  to  $R_B$
    4. add  $(e_i, c_i)$  to  $R_C$



# Converting Non-Binary Relationships

- Also need to translate constraints
  - Translating all constraints may not be possible
  - There may be instances in the translated schema that cannot correspond to any instance of R
    - Exercise: add constraints to the relationships  $R_A$ ,  $R_B$  and  $R_C$  to ensure that a newly created entity corresponds to exactly one entity in each of entity sets A, B and C
  - We can avoid creating an identifying attribute by making E a weak entity set (described shortly) identified by the three relationship sets

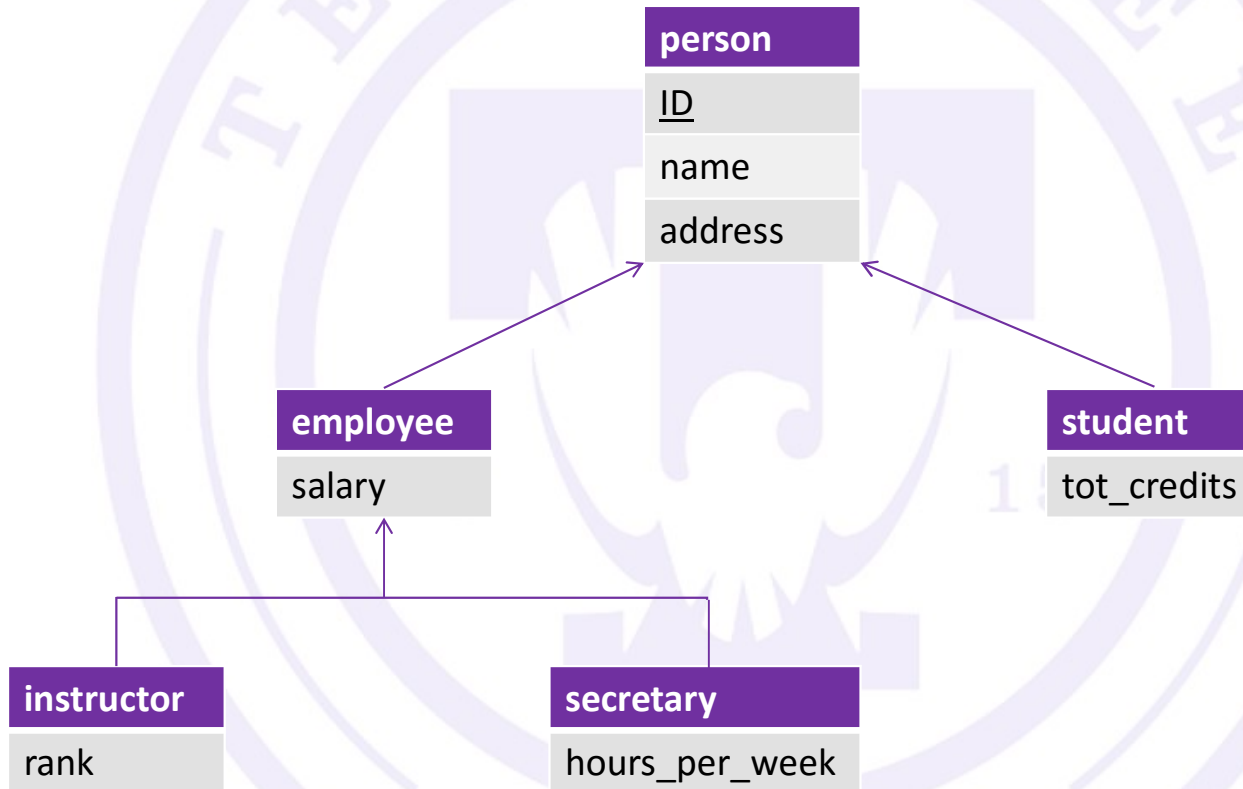


# Extended E-R Features : Specialization

- Top-down design process; we designate subgroupings within an entity set that are distinctive from other entities in the set.
- These subgroupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set.
- Depicted by a triangle component labeled ISA (E.g., instructor “is a” person).
- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked.



# Specialization Example



# Extended ER Features : Generalization

- A bottom-up design process – combine a number of entity sets that share the same features into a higher-level entity set.
- Specialization and generalization are simple inversions of each other; they are represented in an E-R diagram in the same way.
- The terms specialization and generalization are used interchangeably.

# Specialization and Generalization

- Can have multiple specializations of an entity set based on different features.
- E.g., permanent\_employee vs. temporary\_employee, in addition to instructor vs. secretary
- Each particular employee would be
  - A member of one of permanent\_employee or temporary\_employee,
  - And also a member of one of instructor, secretary
- The ISA relationship also referred to as **superclass - subclass** relationship

# Design Constraints on Specialization and Generalization

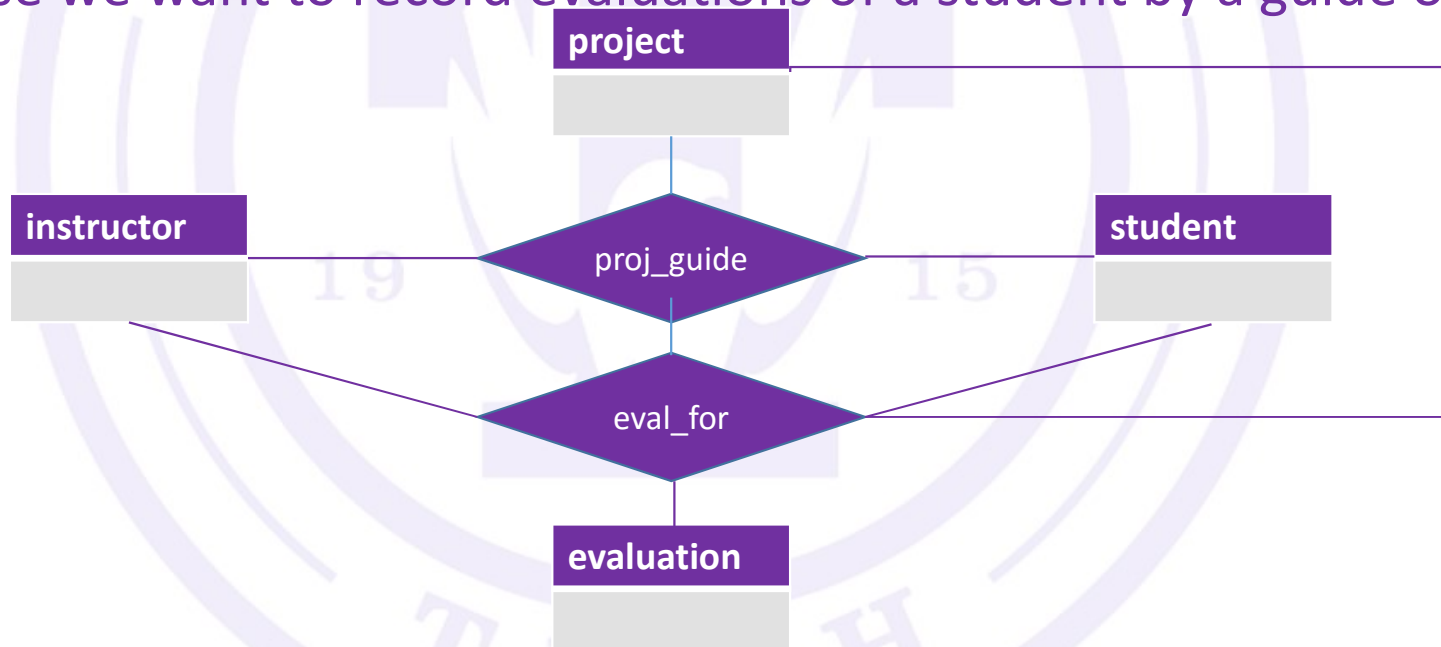
- Constraint on which entities can be members of a given lower-level entity set.
  - Condition-defined
    - Example: all customers over 65 years are members of senior-citizen entity set; senior-citizen ISA person.
  - User-defined
- Constraint on whether or not entities may belong to more than one lower-level entity set within a single generalization.
  - **Disjoint**
    - An entity can belong to only one lower-level entity set
    - Noted in E-R diagram by having multiple lower-level entity sets link to the same triangle
  - **Overlapping**
    - An entity can belong to more than one lower-level entity set

# Design Constraints on Specialization and Generalization

- **Completeness constraint** – specifies whether or not an entity in the higher-level entity set must belong to at least one of the lower-level entity sets within a generalization.
  - **Total**: an entity must belong to one of the lower-level entity sets
  - **Partial**: an entity need not belong to one of the lower-level entity sets

# Aggregation

- Consider the ternary relationship proj\_guide, which we saw earlier
- Suppose we want to record evaluations of a student by a guide on a project

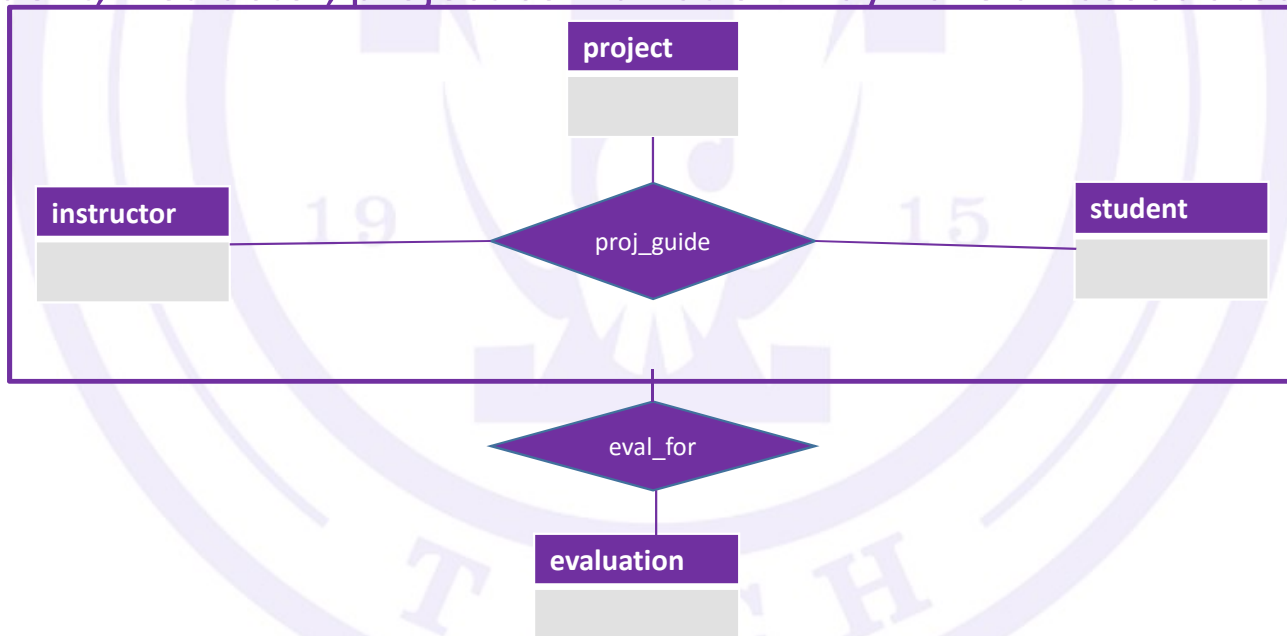


# Aggregation

- Relationship sets `eval_for` and `proj_guide` represent overlapping information
  - Every `eval_for` relationship corresponds to a `proj_guide` relationship
  - However, some `proj_guide` relationships may not correspond to any `eval_for` relationships
    - So we can't discard the `proj_guide` relationship
- Eliminate this redundancy via aggregation
  - Treat relationship as an abstract entity
  - Allows relationships between relationships
  - Abstraction of relationship into new entity

# Aggregation

- Without introducing redundancy, the following diagram represents:
  - A student is guided by a particular instructor on a particular project
  - A student, instructor, project combination may have an associated evaluation





# Representing Specialization via Schemas

- Method 1:

- Form a schema for the higher-level entity
- Form a schema for each lower-level entity set, include primary key of higher-level entity set and local attributes

schema	attributes
person	ID, name, street, city
student	ID, name, street, city, tot_cred
employee	ID, name, street, city, salary

- Drawback: getting information about, an *employee* requires accessing two relations, the one corresponding to the low-level schema and the one corresponding to the high-level schema

# Representing Specialization as Schemas

- Method 2:
  - Form a schema for each entity set with all local and inherited attributes

schema	attributes
person	ID, name, street, city
student	ID, tot_cred
employee	ID, salary

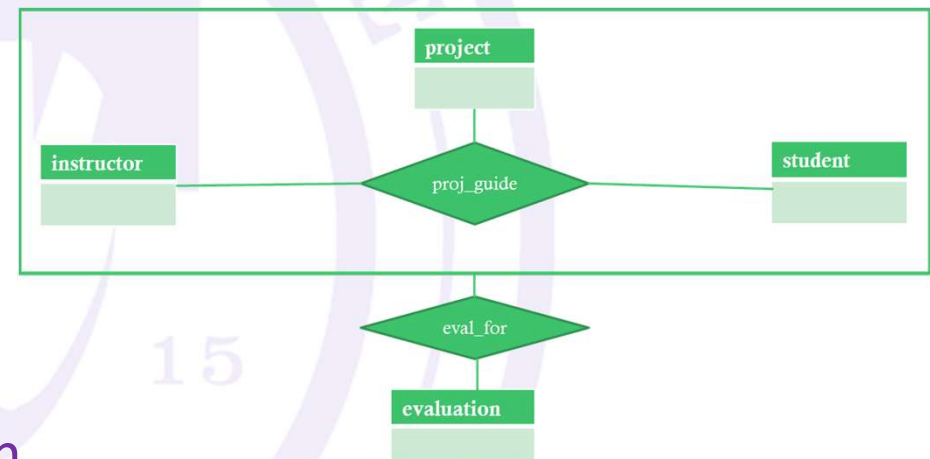
- If specialization is total, the schema for the generalized entity set (*person*) not required to store information
  - Can be defined as a “view” relation containing union of specialization relations
  - But explicit schema may still be needed for foreign key constraints
- Drawback: *name*, *street* and *city* may be stored redundantly for people who are both students and employees

# Schemas Corresponding to Aggregation

- To represent aggregation, create a schema containing
  - Primary key of the aggregated relationship,
  - The primary key of the associated entity set
  - Any descriptive attributes

## Schemas Corresponding to Aggregation

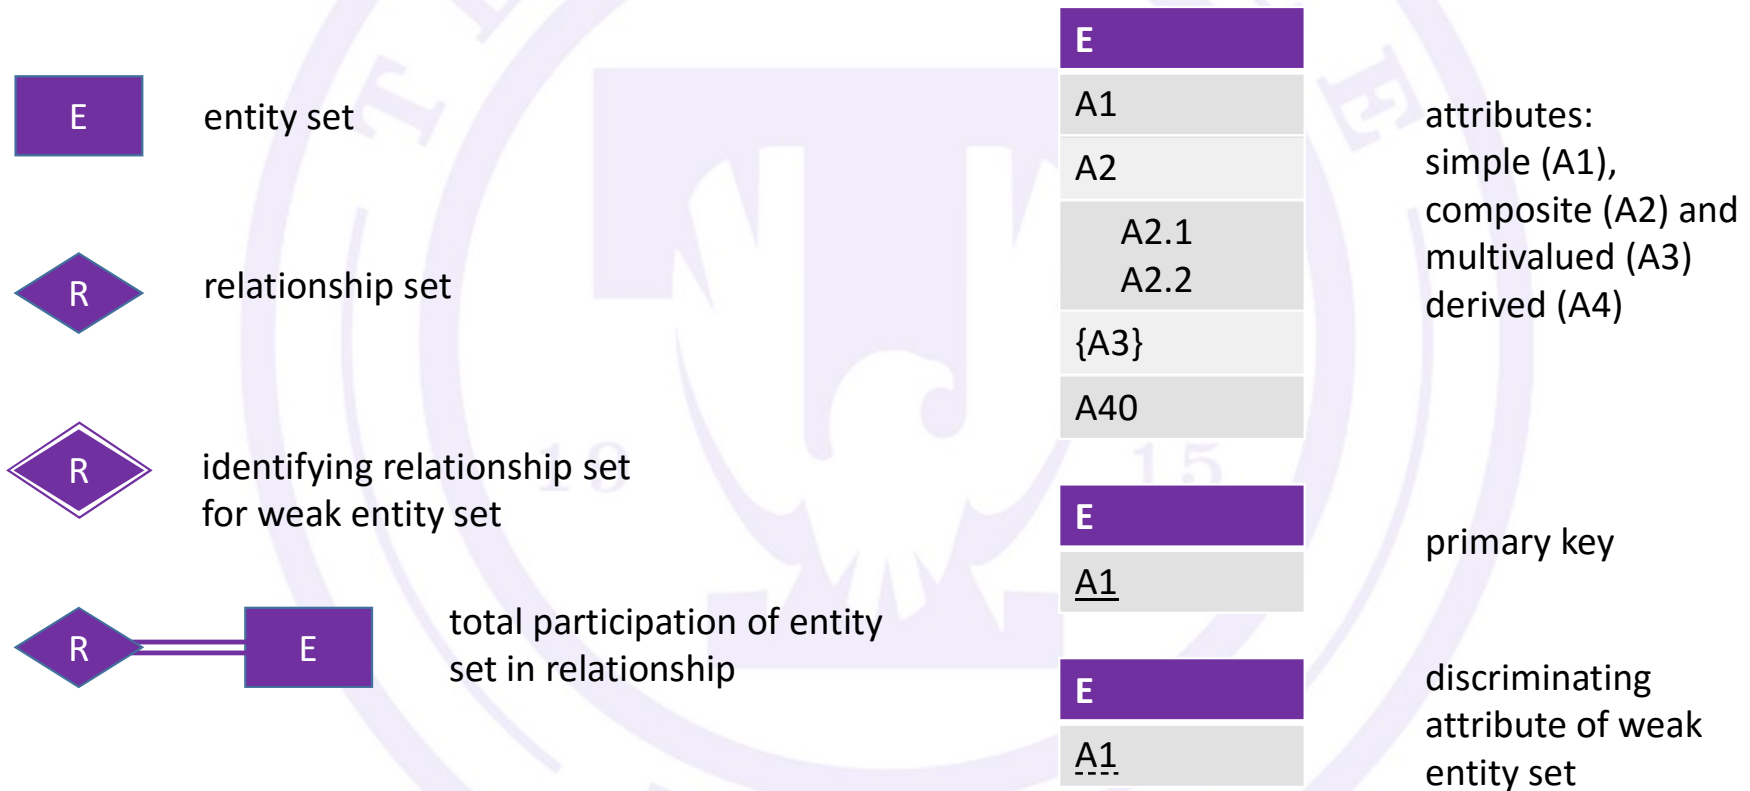
- For example, to represent aggregation manages between relationship works\_on and entity set manager, create a schema  
eval\_for (s\_ID, project\_id, i\_ID, evaluation\_id)
- Schema proj\_guide is redundant provided we are willing to store null values for attribute manager\_name in relation on schema manages



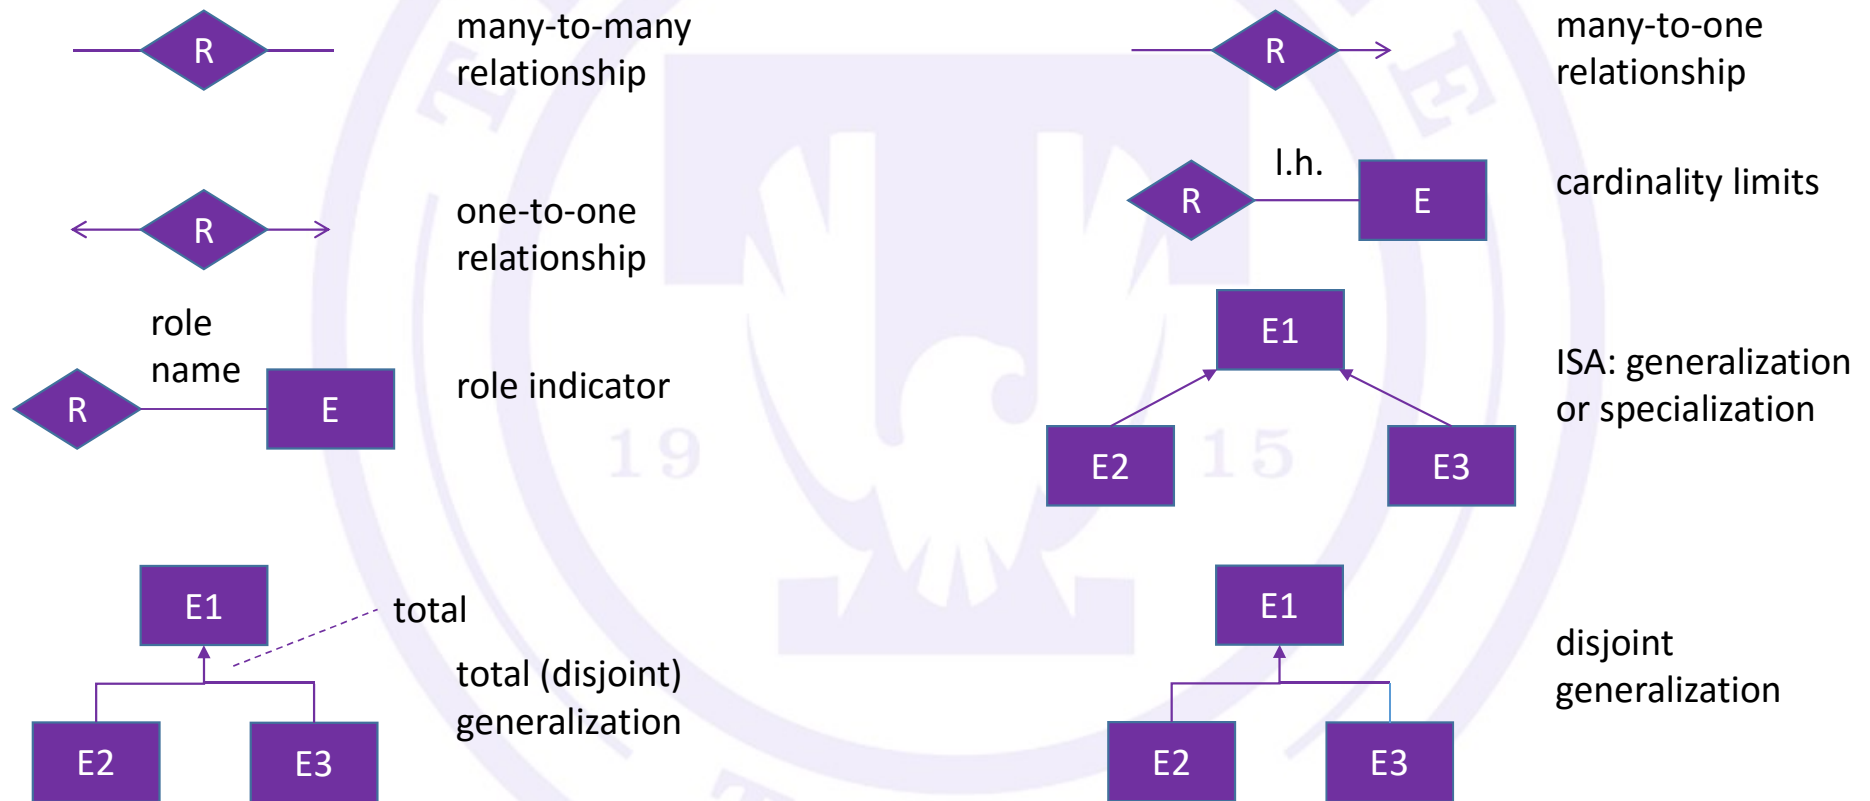
# E-R Design Decisions

- The use of an attribute or entity set to represent an object.
- Whether a real-world concept is best expressed by an entity set or a relationship set.
- The use of a ternary relationship versus a pair of binary relationships.
- The use of a strong or weak entity set.
- The use of specialization/generalization – contributes to modularity in the design.
- The use of aggregation – can treat the aggregate entity set as a single unit without concern for the details of its internal structure.

# Summary of Symbols Used in E-R Notation




# Summary of Symbols Used in E-R Notation



# Alternative ER Notations

- Chen, IDE1FX, ...

entity set E with simple attribute A1, composite attribute A2, multivalued attribute A3, derived attribute A4 and primary key A1



weak entity set



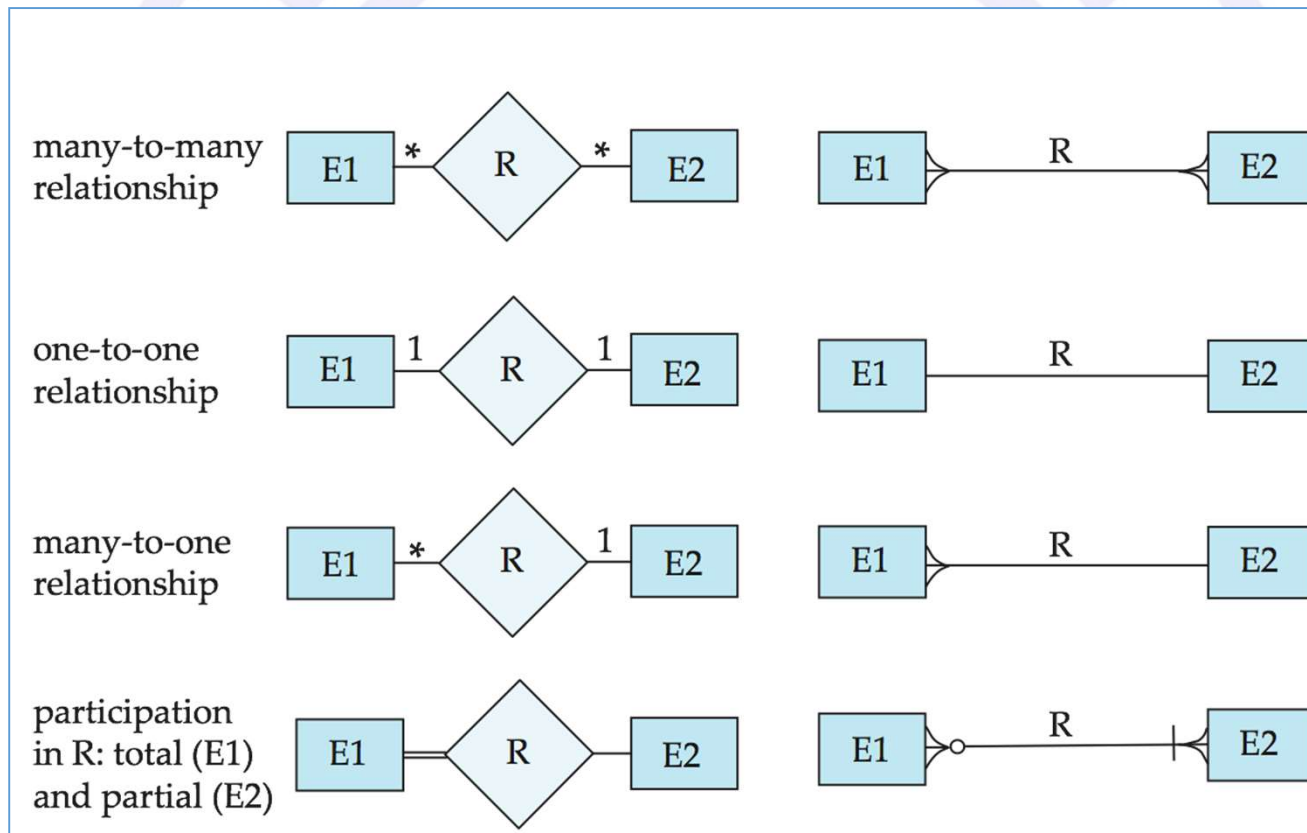
generalization



total  
generalization



# Alternative ER Notations


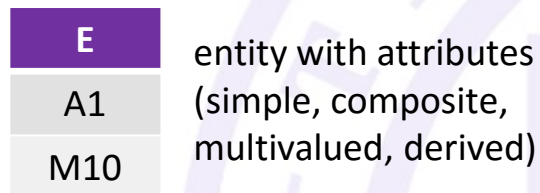


# UML


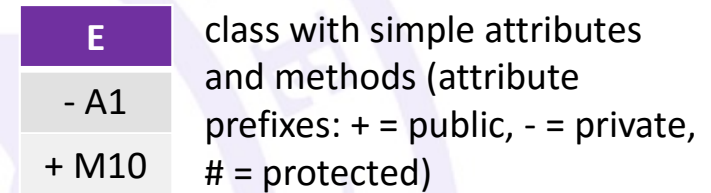
- **UML**: Unified Modeling Language
  - UML has many components to graphically model different aspects of an entire software system
  - UML Class Diagrams correspond to E-R Diagram, but several differences.

# ER v UML Class Diagrams

## ER Diagram Notation



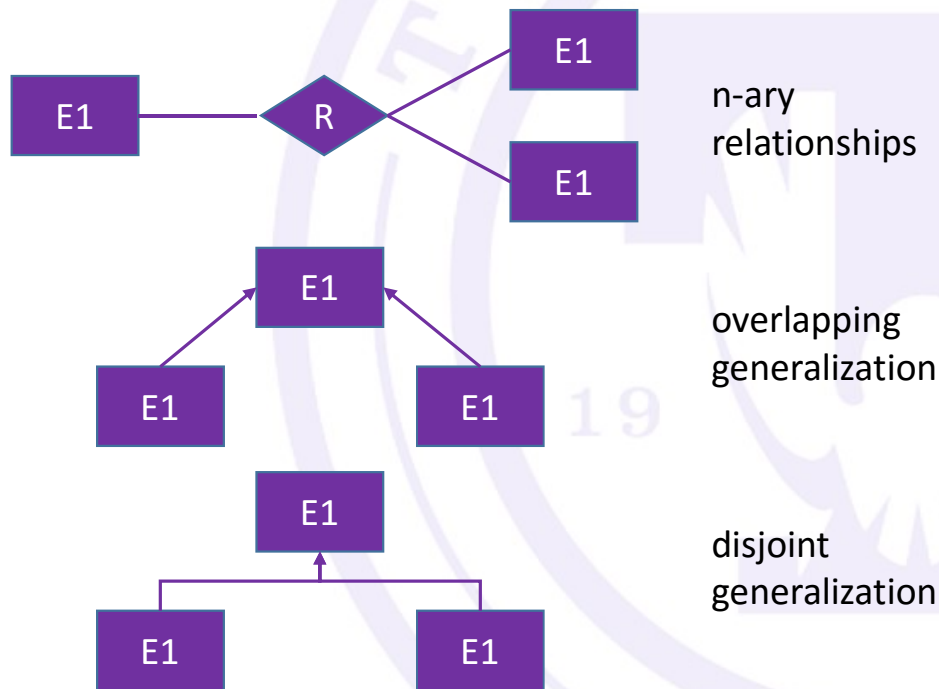
## Equivalent in UML




\* Note the reversal of position in cardinality constraints depiction

# ER v UML Class Diagrams

ER Diagram Notation



Equivalent in UML



\*Generalization can use merged or separate arrows independent of disjoint/overlapping

# ER v UML Class Diagrams

- Binary relationship sets are represented in UML by just drawing a line connecting the entity sets. The relationship set name is written adjacent to the line.
- The role played by an entity set in a relationship set may also be specified by writing the role name on the line, adjacent to the entity set.
- The relationship set name may alternatively be written in a box, along with attributes of the relationship set, and the box is connected, using a dotted line, to the line depicting the relationship set.

End of Chapter 7

