

Chapter 5

Advanced SQL

Slides by Silberschatz, Modifications by Rogers and Brown

Procedural Extensions and Stored Procedures

- SQL Provides a **module** language
 - Permits definition of procedures in SQL, with if-then-else statements, for and while loops, etc
- Stored Procedures
 - Can store procedures in the database
 - Then execute them using the **call** statement
 - Permit external applications to operate on the database without knowing about internal details
- Object-oriented aspects of these features are covered in Chapter 22 (Object Oriented Databases)
 - We will not get THAT far this term

Functions and Procedures

- SQL:1999 supports functions and procedures
 - Functions/Procedures can be written in SQL itself, or in an external programming language
 - Functions are particularly useful with specialized data types such as images and geometric objects
 - Example: Functions to check if polygons overlap or to compare images for similarity
 - Some database systems support **table-valued functions**, which can return a relation as a result
- SQL:1999 also supports a rich set of imperative constructs, including :
 - Loops
 - If-Then-Else
 - Assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999

SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department


```
create function dept_count (dept_name varchar (20))
returns integer
begin
    declare d_count integer;
    select count ( * ) into d_count
    from instructor
    where instructor.dept_name = dept_name
    return d_count;
end
```
- Find the department name and budget of all departments with more than 12 instructors


```
select dept_name, budget
from department
where dept_count (dept_name) > 1
```

Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all accounts owned by a given customer

```
create function instructors_of (dept_name char (20)
returns table (ID          varchar (5),
               name        varchar (20),
               dept_name    varchar (20),
               salary       numeric (8, 2))
return table
( select ID, name, dept_name, salary
  from instructor
  where instructor.dept_name =
        instructors_of.dept_name)
```

- Usage

```
select *
from table (instructors_of ('Music'))
```

SQL Procedures

- The dept_count function could instead be written as a procedure:

```
create procedure dept_count_proc (in dept_name varchar (20),
                                out d_count integer)
begin
    select count ( * ) into d_count
    from instructor
    where instructor.dept_name = dept_count_proc.dept_name
end
```

- Procedures can be invoked either from an SQL procedure or from embedded SQL using the call statement

```
declare d_count integer;
call dept_count_proc ('Physics', d_count);
```

Procedures and functions can be invoked also from dynamic SQL

- SQL:1999 allows more than one function/procedure of the same name (called **overloading**), as long as the number of arguments differ or at least the types of arguments differ

Procedural Constructs

- Warning: most database systems implement their own variant of the standard syntax
 - Read the system manual for the system you are using for any variations
- Compound statement: begin ... end,
 - May contain multiple SQL statements between begin and end
 - Local variables can be declared within a compound statements
- While and repeat statements:

```

declare n integer default 0;
while n < 10 do
    set n = n + 1
end while

repeat
    set n = n - 1
until n = 0
end repeat

```

Procedural Constructs

- For loop
 - Permits iteration over all results of a query
 - Example:

```

declare n integer default 0;
for r as
    select budget from department
    where dept_name = 'Music'
do
    set n = n + r.budget
end for

```

Procedural Constructs

- Conditional Statements (**if-then-else**)
SQL:1999 also supports **case** statements similar to C++ case statement
- Example procedure: registers a student after ensuring classroom capacity is not exceeded
 - Returns a 0 on success, and -1 if capacity would be exceeded
 - See the book for more details
- Signaling of exception conditions, and declaring handlers for exceptions


```

declare out_of_classroom_seats condition
declare exit handler for out_of_classroom_seats
begin
    ...
..signal out_of_classroom_seats
end
      
```

 - The handler here is **exit** – causes enclosing **begin..end** statements to be exited
 - Other actions possible on exception

External Language Functions/Procedures

- SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions

```

create procedure dept_count_proc (in dept_name varchar (20),
                                out count integer)
language C
external name '/usr/avi/bin/dept_count_proc'

create function dept_count (dept_name varchar (20) )
returns integer
language C
external name '/usr/avi/bin/dept_count'
      
```

External Language Functions/Procedures

- Benefits of external language functions/procedures:
 - More efficient for many operations, and more expressive power.
- Drawbacks
 - Code to implement function may need to be loaded into database system and executed in the database system's address space.
 - Risk of accidental corruption of database structures
 - Security risk, allowing users access to unauthorized data
 - There are alternatives, which give good security at the cost of potentially worse performance.
 - Direct execution in the database system's space is used when efficiency is more important than security.

Security with External Language Routines

- To deal with security problems
 - Use **sandbox** techniques
 - That is, use a safe language like Java, which cannot be used to access/damage other parts of the database code.
 - Or, run external language functions/procedures in a separate process, with no access to the database process' memory.
 - Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space.

Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - **Note** : Syntax illustrated here may not work exactly on your database system; check the system manuals

Trigger Example

- E.g. time_slot_id is not a primary key of timeslot, so we cannot create a foreign key constraint from section to timeslot
- Alternative: use triggers on section and timeslot to enforce integrity constraints

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from time_slot)) /* time_slot_id not present in time_slot */
begin
    rollback
end;
```

Trigger Example (cont)

- create trigger timeslot_check2 after delete on timeslot
referencing old row as orow
for each row
when (orow.time_slot_id not in (
 select time_slot_id
 from time_slot)
/* last tuple for time_slot_id deleted from time_slot */
and orow.time_slot_id in (
 select time_slot_id
 /* and time_slot_id still referenced from section */
 from section))
begin
 rollback
end;

Triggering Events and Actions in SQL

- Triggering event can be insert, delete or update
- Triggers on update can be restricted to specific attributes
 - E.g., after update of takes on grade
- Values of attributes before and after an update can be referenced
 - referencing old row as : for deletes and updates
 - referencing new row as : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints.
E.g. convert blank grades to null.

```
create trigger setnull_trigger before update of takes
referencing new row as nrow
for each row
when (nrow.grade = ' ')
begin atomic
    set nrow.grade = null;
end;
```


Trigger to Maintain credits_earned value

- create trigger credits_earned after update of takes on (grade)
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
and (orow.grade = 'F' or orow.grade is null)
begin atomic
 update student
 set tot_cred = tot_cred +
 (select credits
 from course
 where course.course_id= nrow.course_id)
 where student.id = nrow.id;
end;

Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - Use for each statement instead of for each row
 - Use referencing old table or referencing new table to refer to temporary tables (called transition tables) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows

When Not to Use Triggers

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
 - Define methods to update fields
 - Carry out actions as part of the update methods instead of through a trigger

When Not to Use Triggers

- Risk of unintended execution of triggers, for example, when
 - Loading data from a backup copy
 - Replicating updates at a remote site
 - Trigger execution can be disabled before such actions.
- Other risks with triggers:
 - Error leading to failure of critical transactions that set off the trigger
 - Cascading execution

End of Chapter 5

